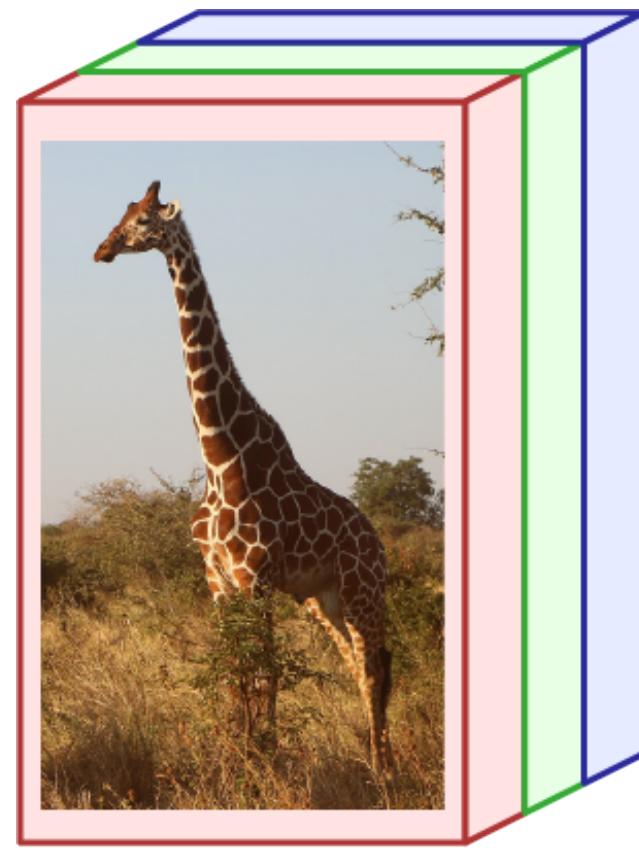
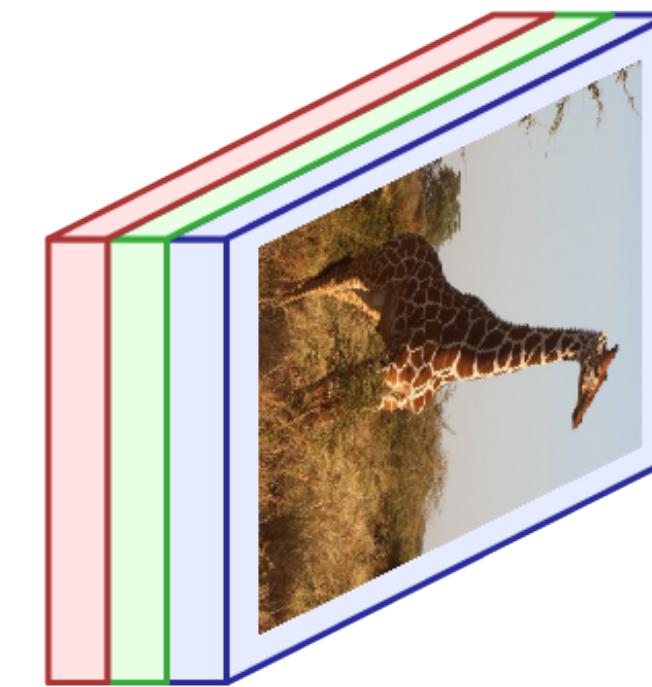


# Images



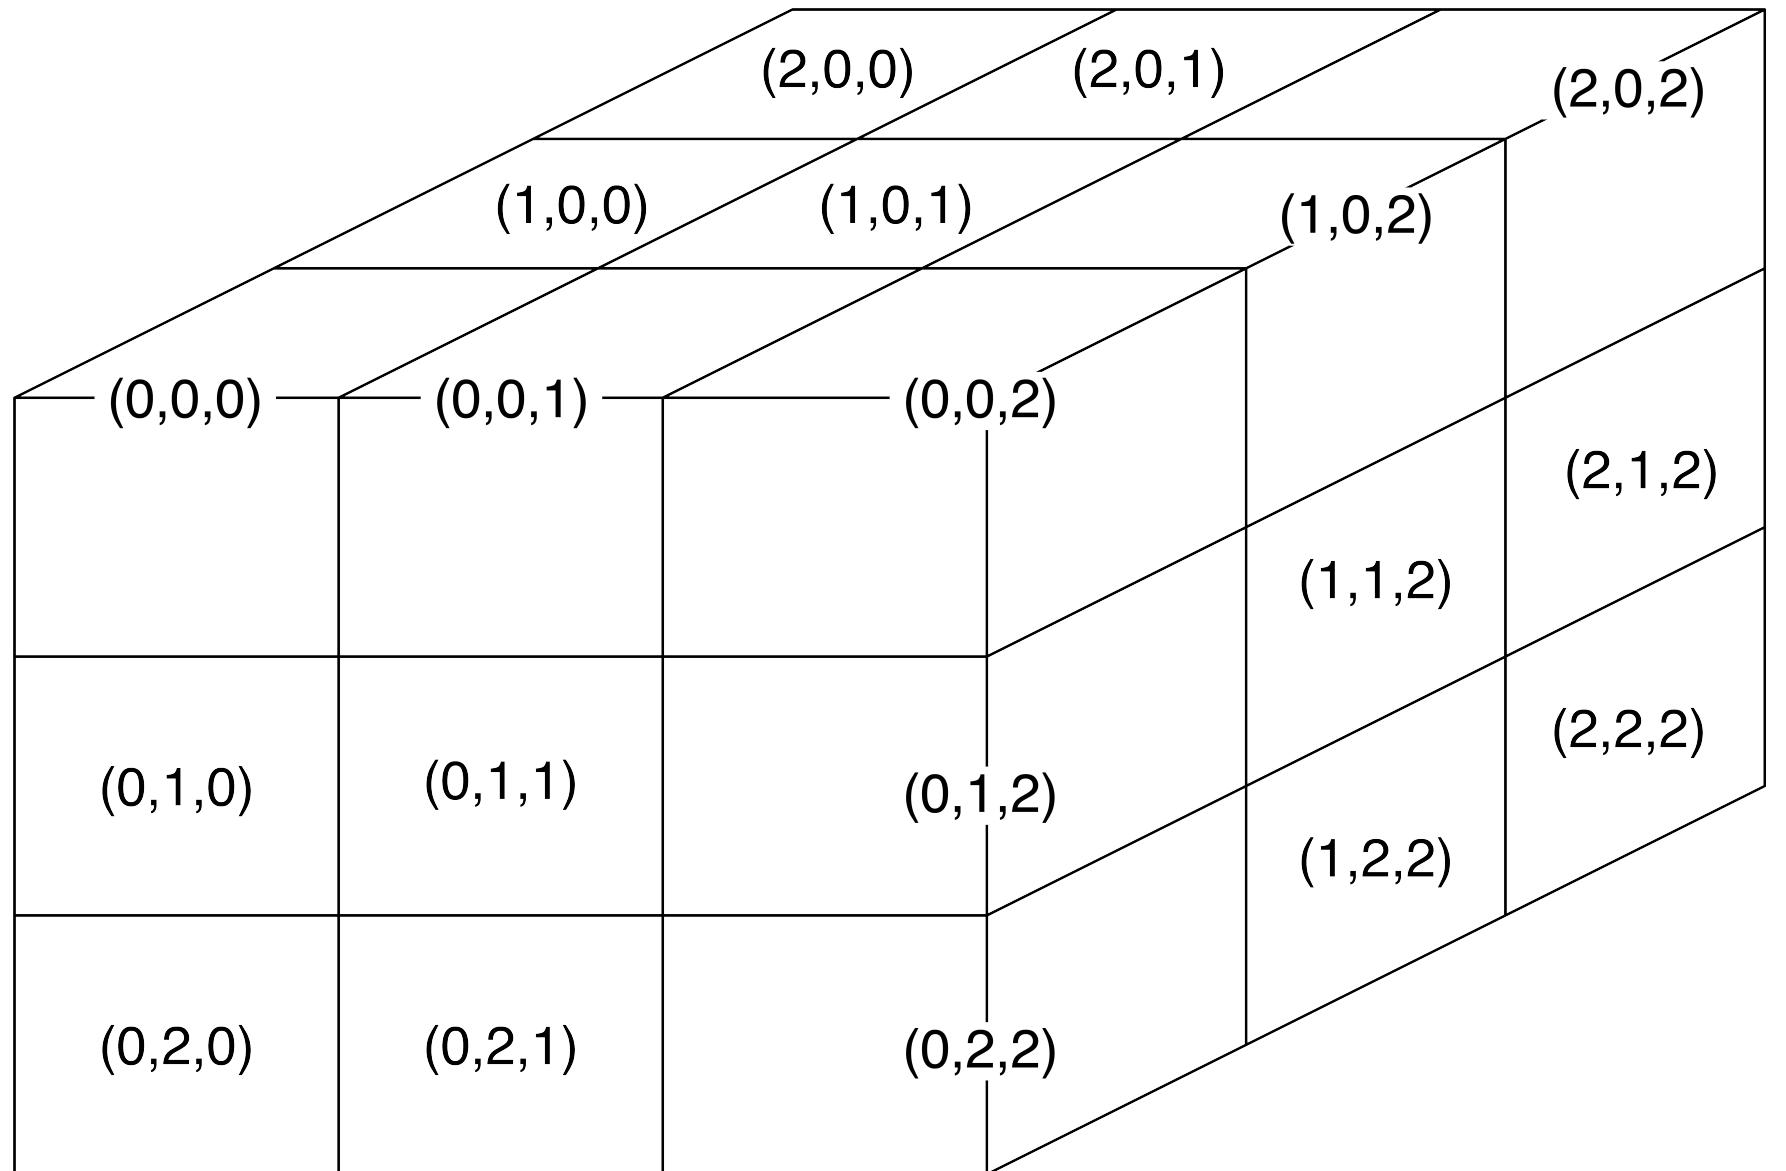
(a)



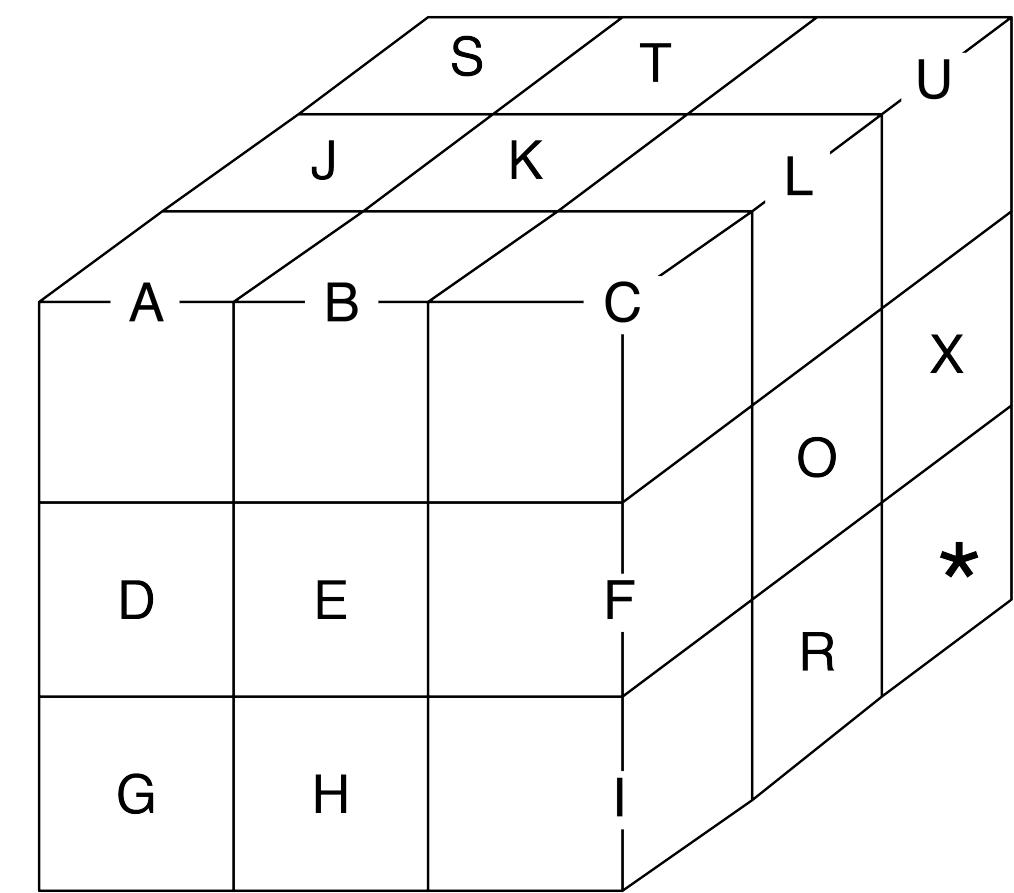
(b)

(a) Channels First (Keras) (b) Channels last (pytorch)

# Channels first arrangement

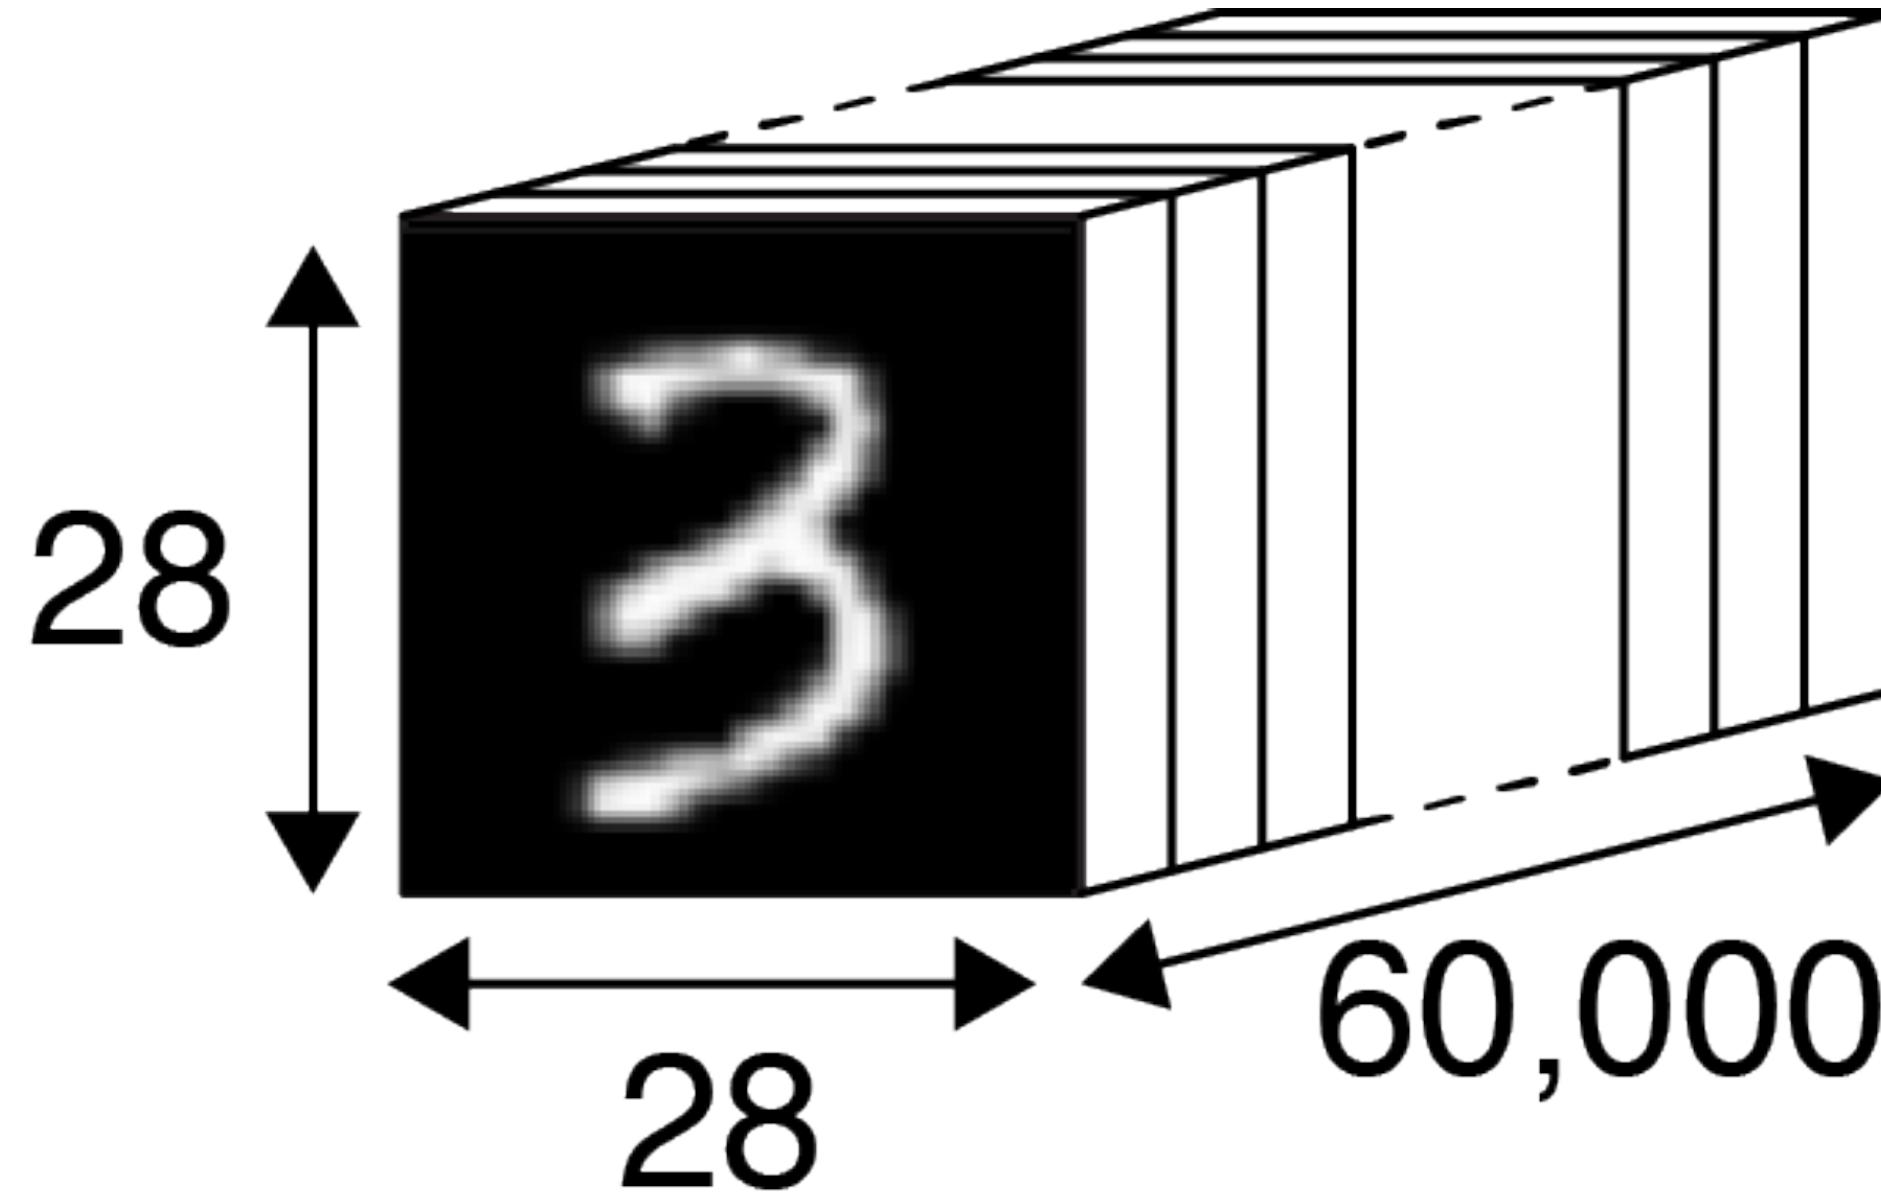


(a)



(b)

# What about multiple images?



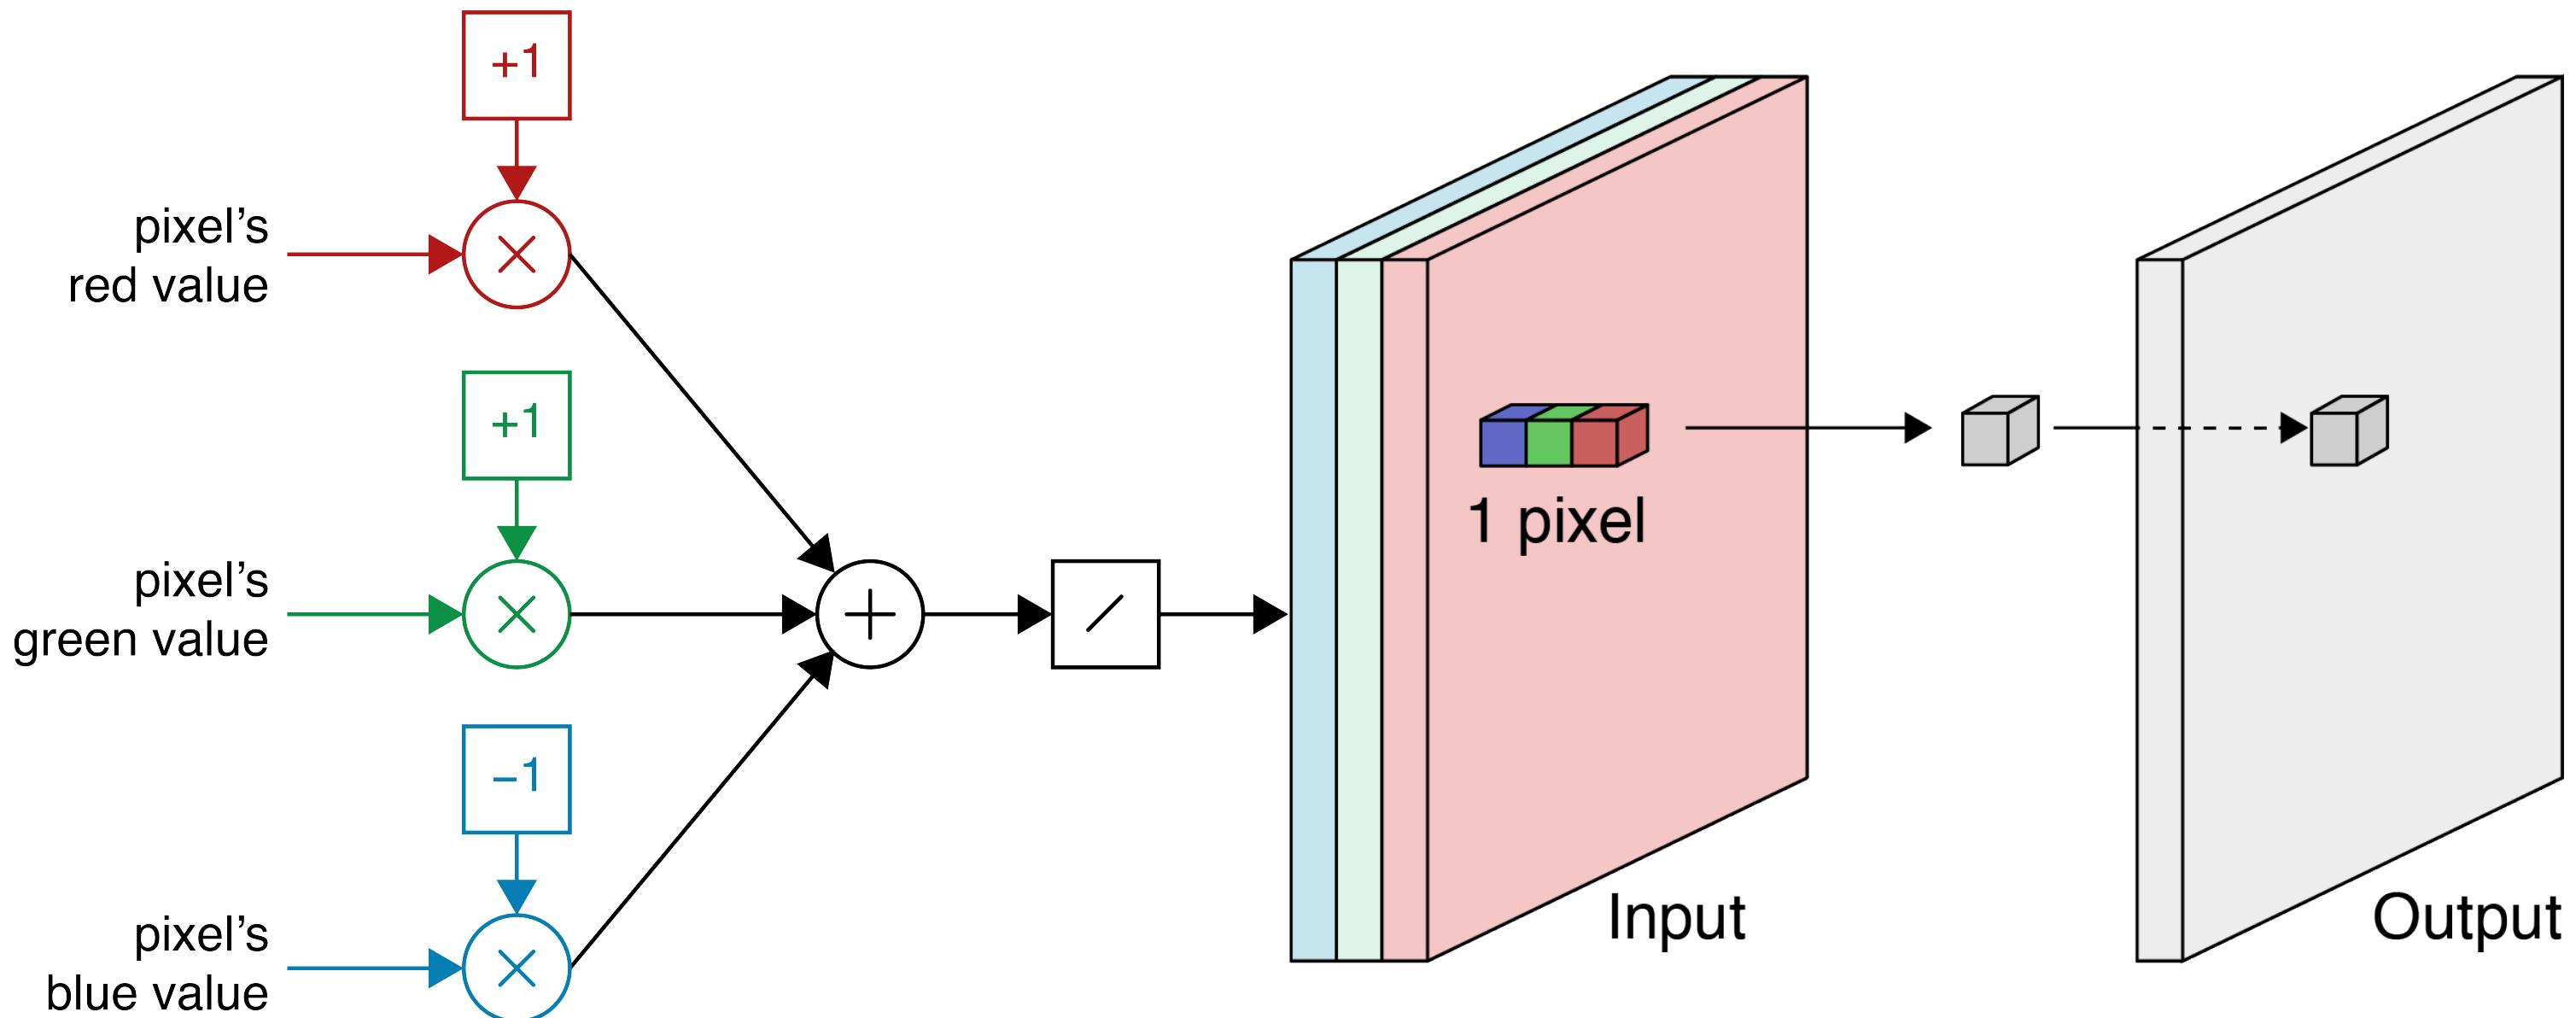
# And a single image?

# Convolutional Networks

- pay attention to the spatial locality of images
- this is done through the use of "filters"
- thus the representations learnt are spatial and bear a mapping to reality
- and are hierarchical..later layers learn features composed from the previous layers
- perhaps even approximating what the visual cortex does.

# Deep Learning Components

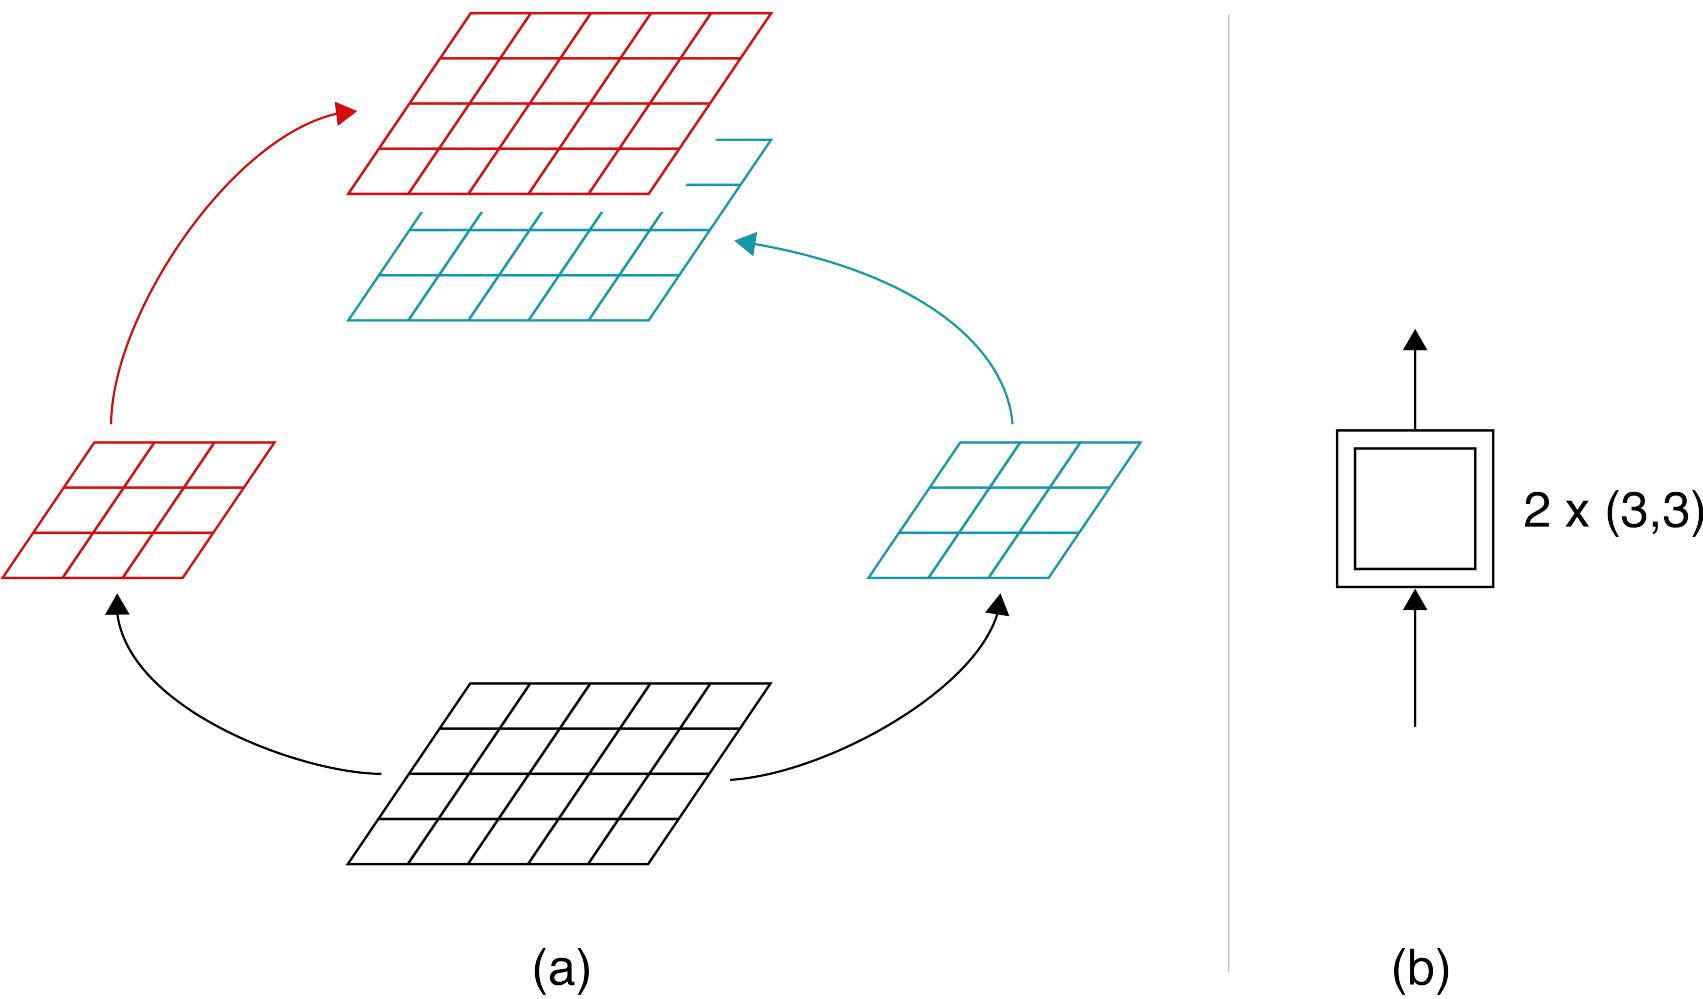
# The idea of a filter: detecting yellow



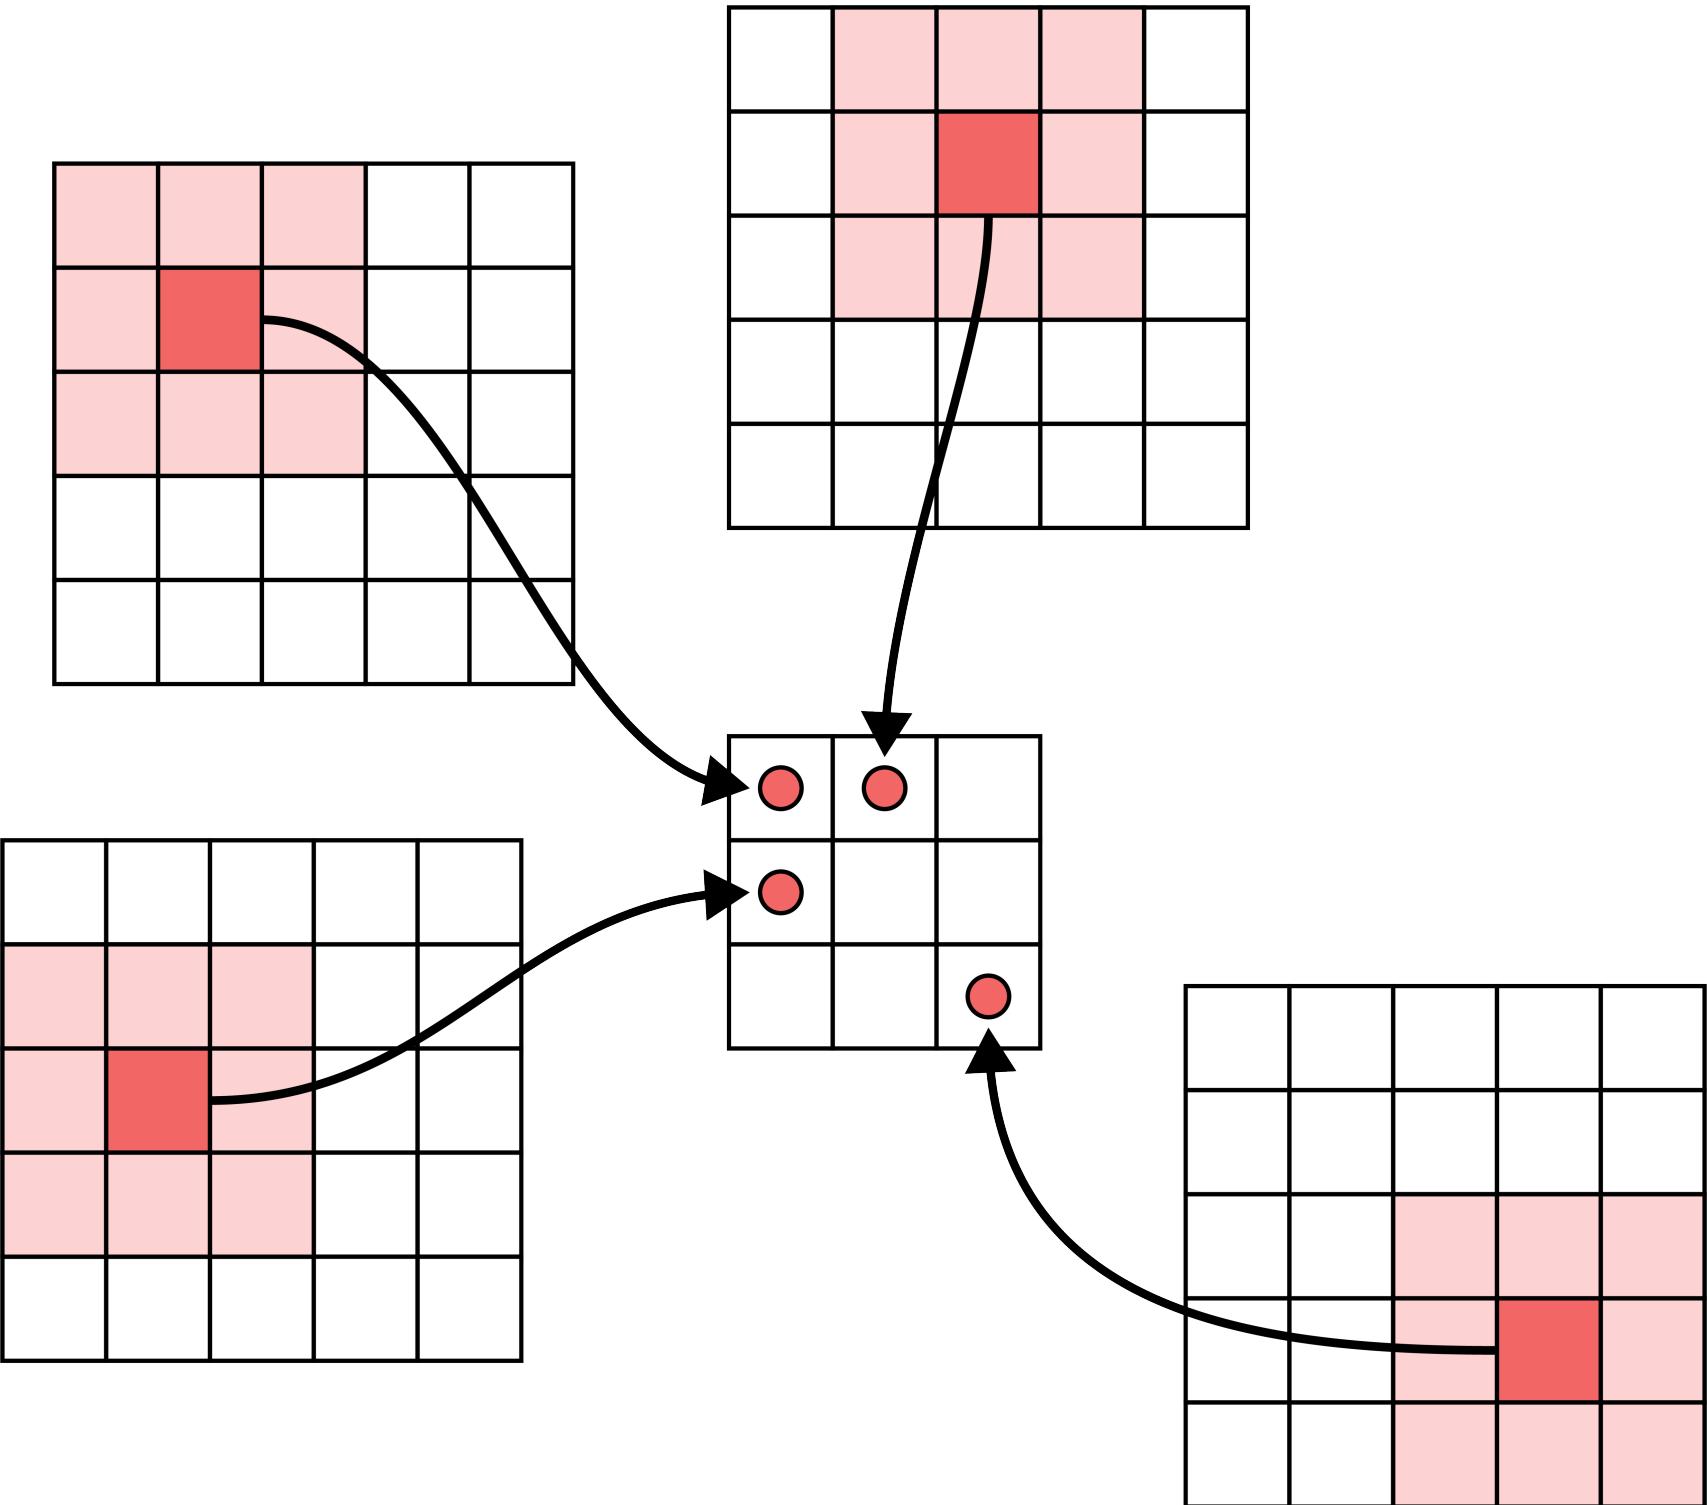


# Convolution Layer

- applies a convolution kernel to an image
- the kernel may have depth, and can use padding on the original image and strides in deciding how to move the kernel window over the image
- the output image(s) depend on these paddings, strides, size of kernel, depth of kernel (channels)
- the kernel usually has weights which multiply the local 'x' in the image and perhaps add bias..in other words, a linear convolution



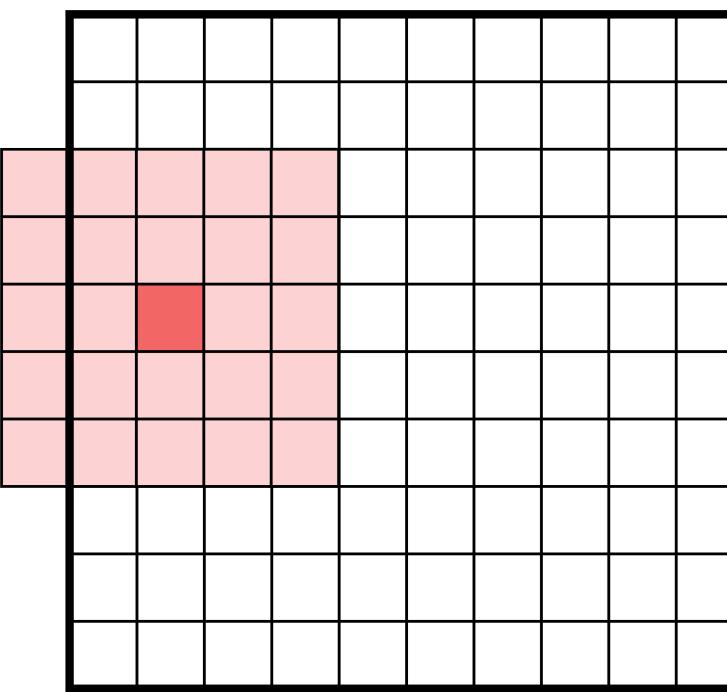
# What is a convolution



- slide 3x3 filter across image of size 5x5
- filters must fit into image, thus output will be of smaller size, in this case 3x3
- multiply values in filter by values of pixel and add a bias:  $(w, b) \cdot (x, 1)$

# Padding

What happens at the edge? We would fall off. Or we can't go to the edge.



- Image size decrease by a convolution is called a "valid" convolution.
- Keeping the same size by 0-padding is called a "same" convolution

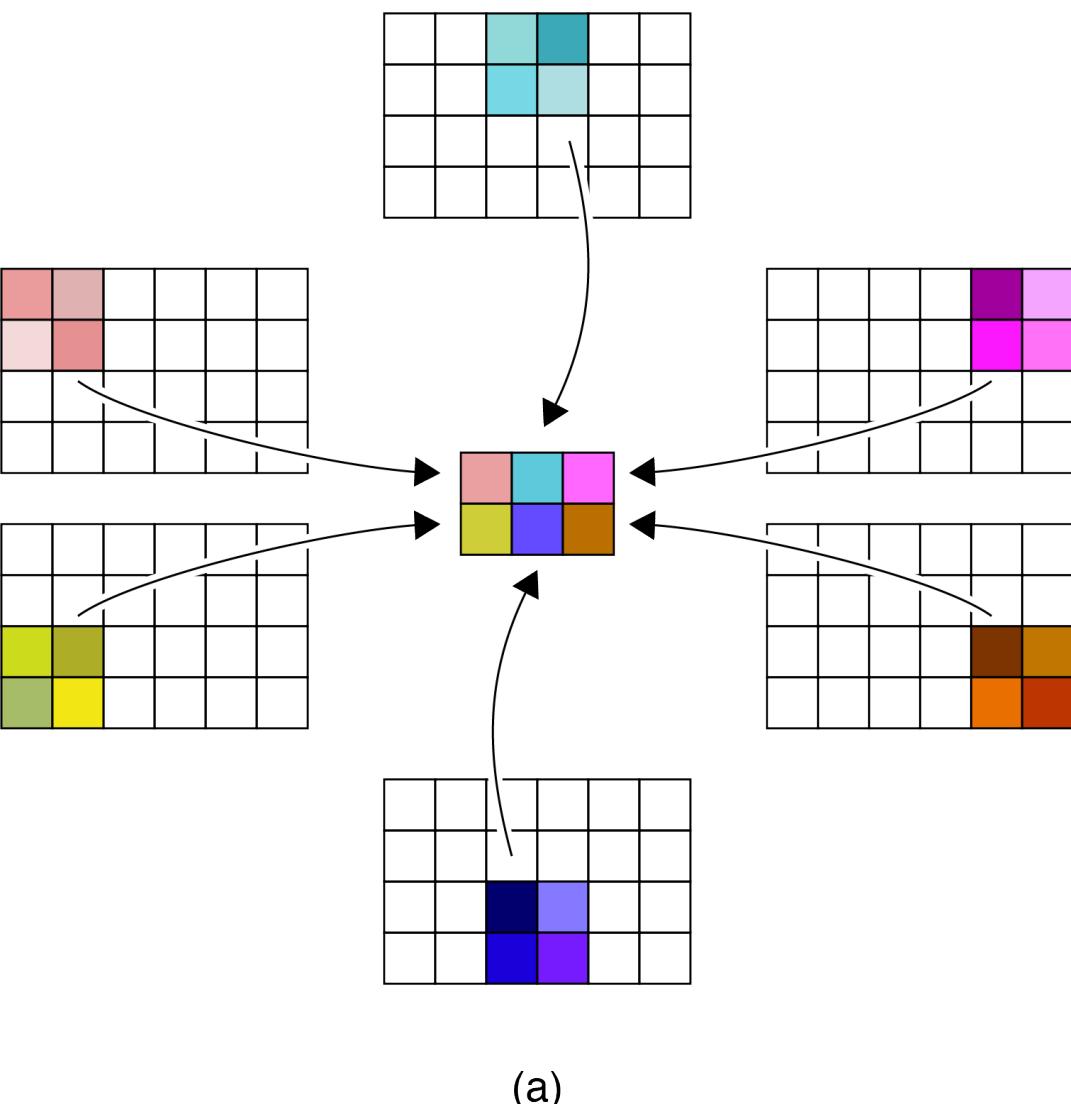
# Downsampling

Eventually we want to get to larger length scales across an image, to learn features such as eyes and ears.

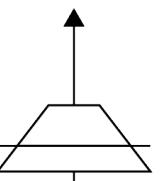
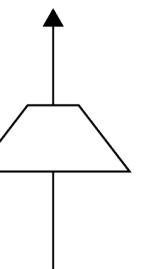
There are two ways to accomplish this:

- pooling, or combining neighboring pixels
- striding, or reducing image size by skipping every nth pixel

# Pooling Layer



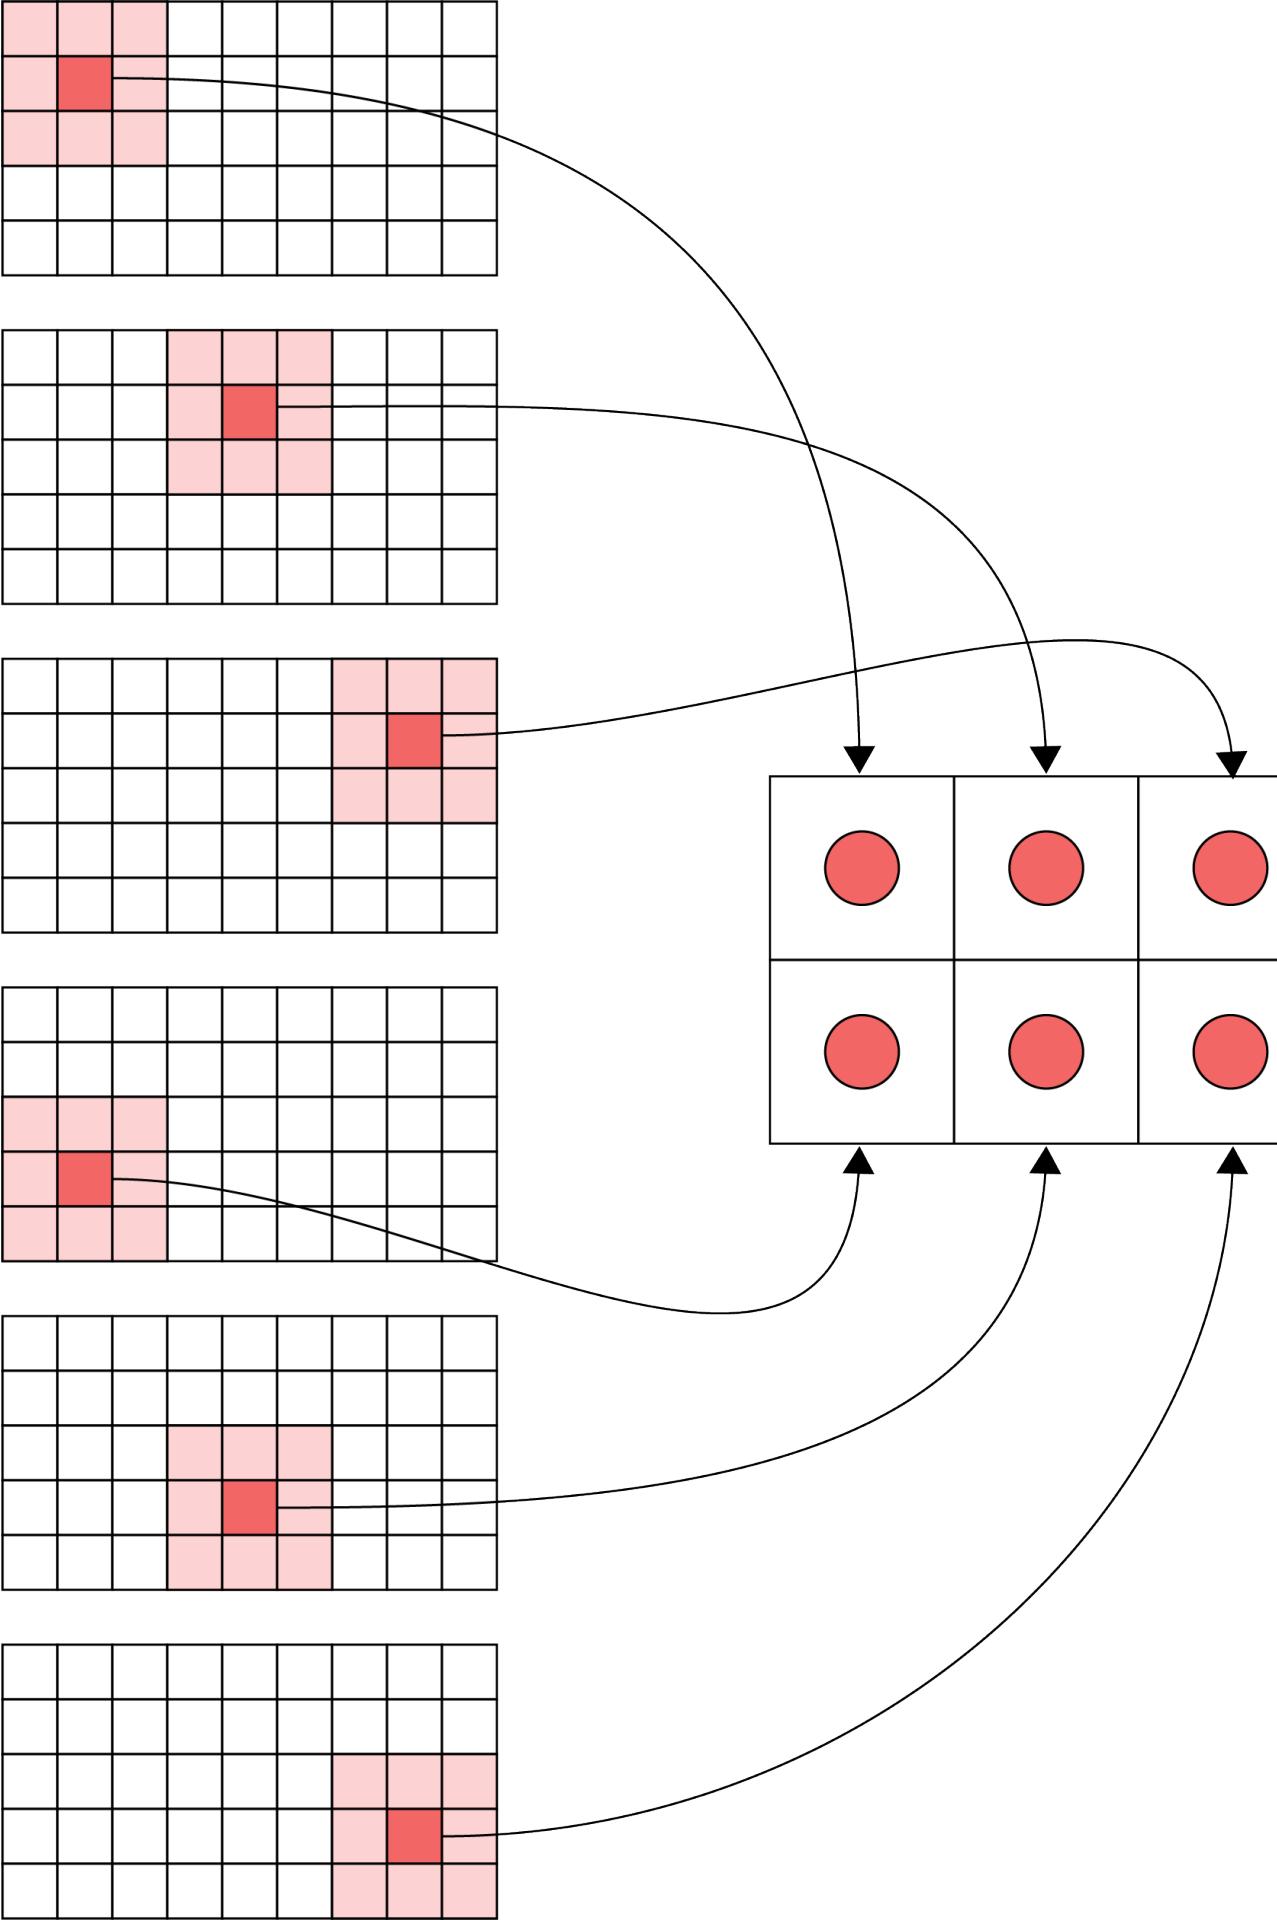
average  
pooling      maximum  
pooling



- down samples images, usually using a non-overlapping stride
- 2 common types, max-pooling and average-pooling. Former better for detecting sharp edges
- makes images smaller for manageability, and also helps in learning at the next level in the hierarchy

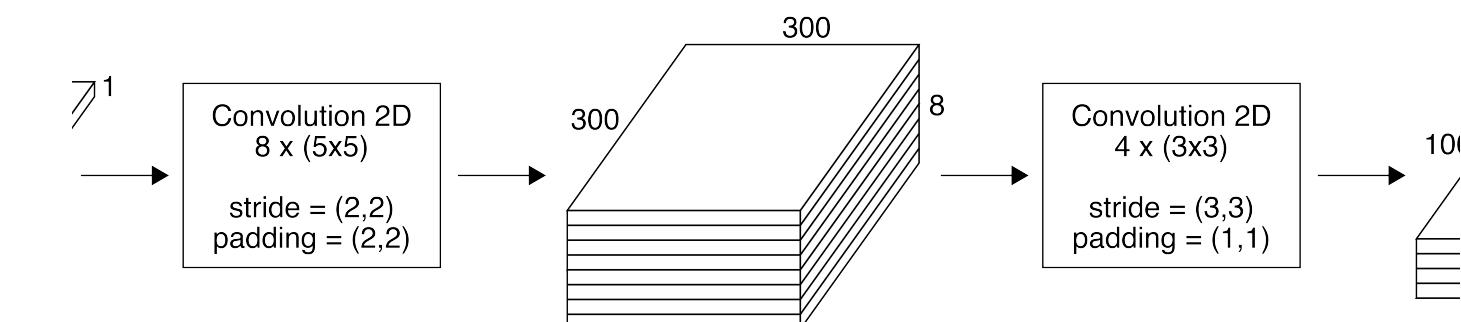
(b)

# Striding

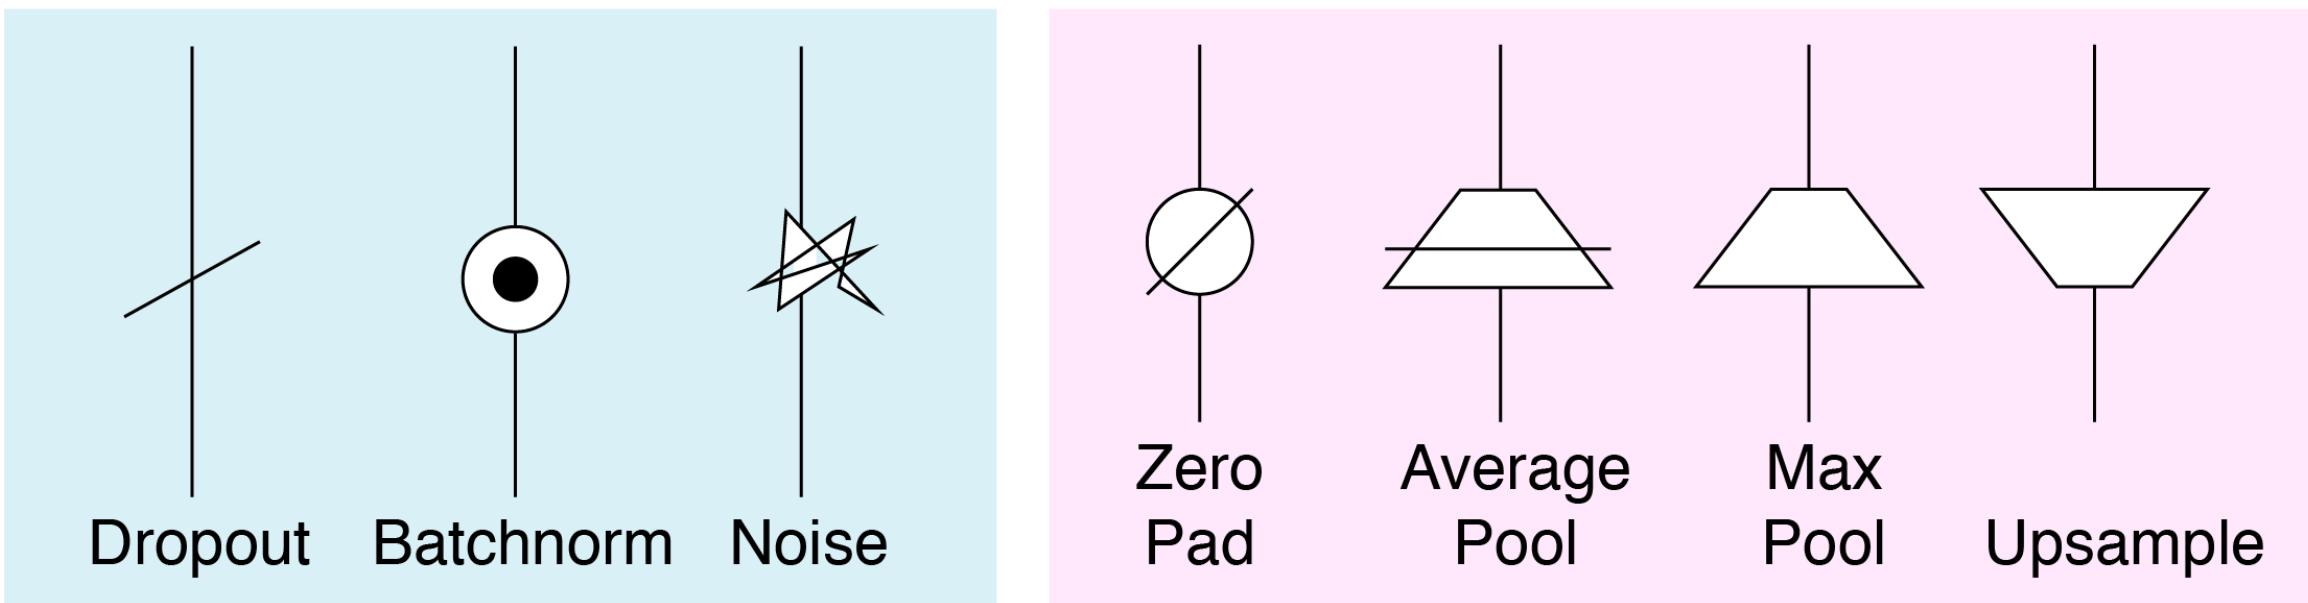
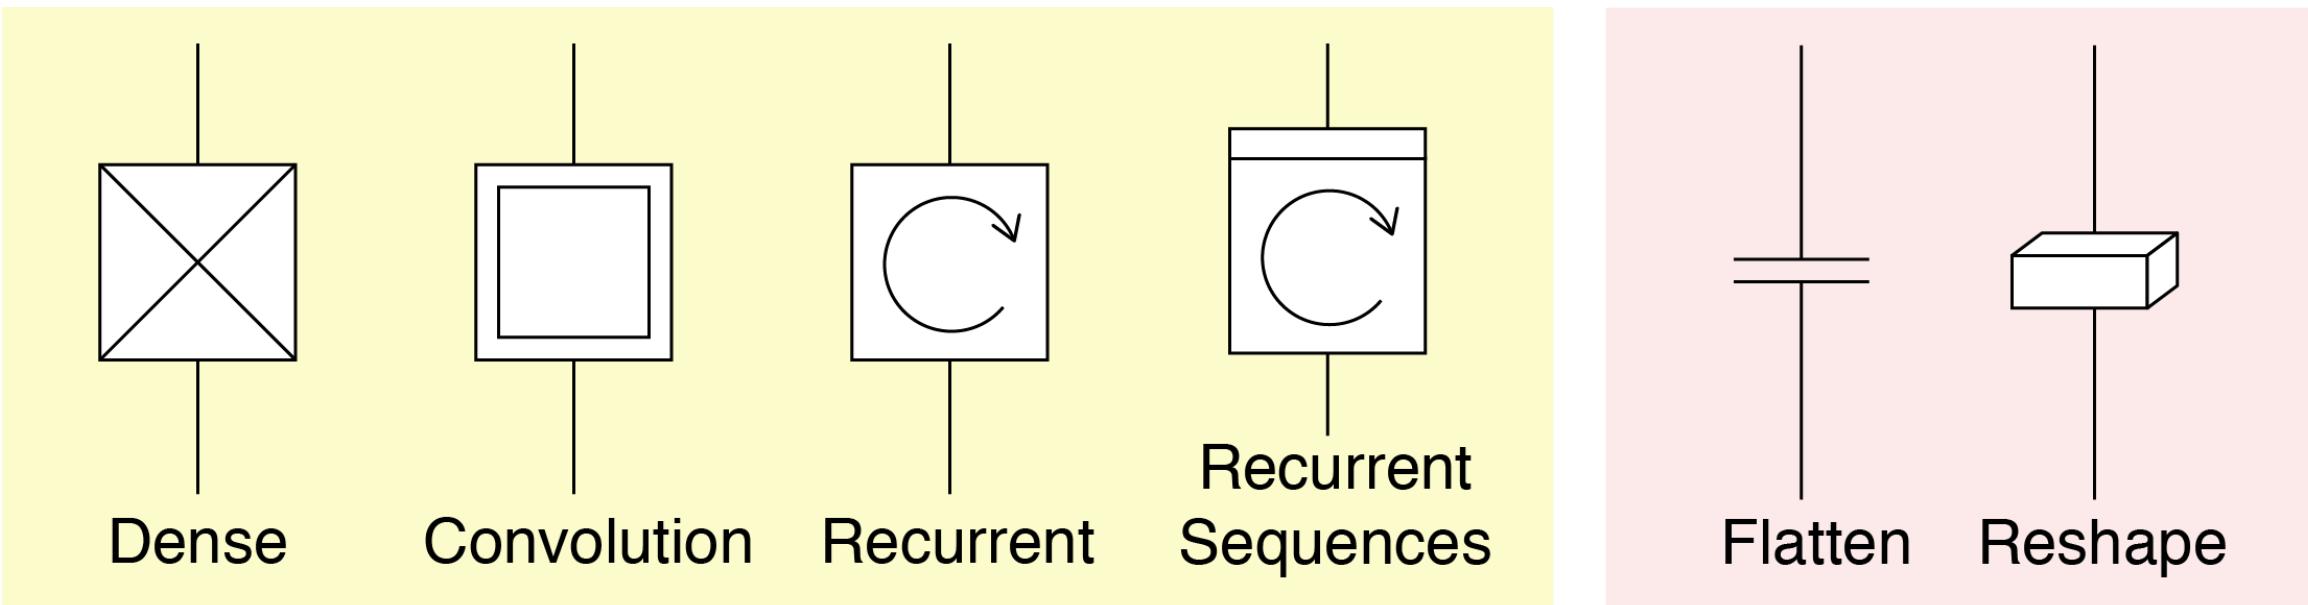


In recent years, experience has shown that striding while convolving is faster and often gives results that are just as good, or better, than pooling

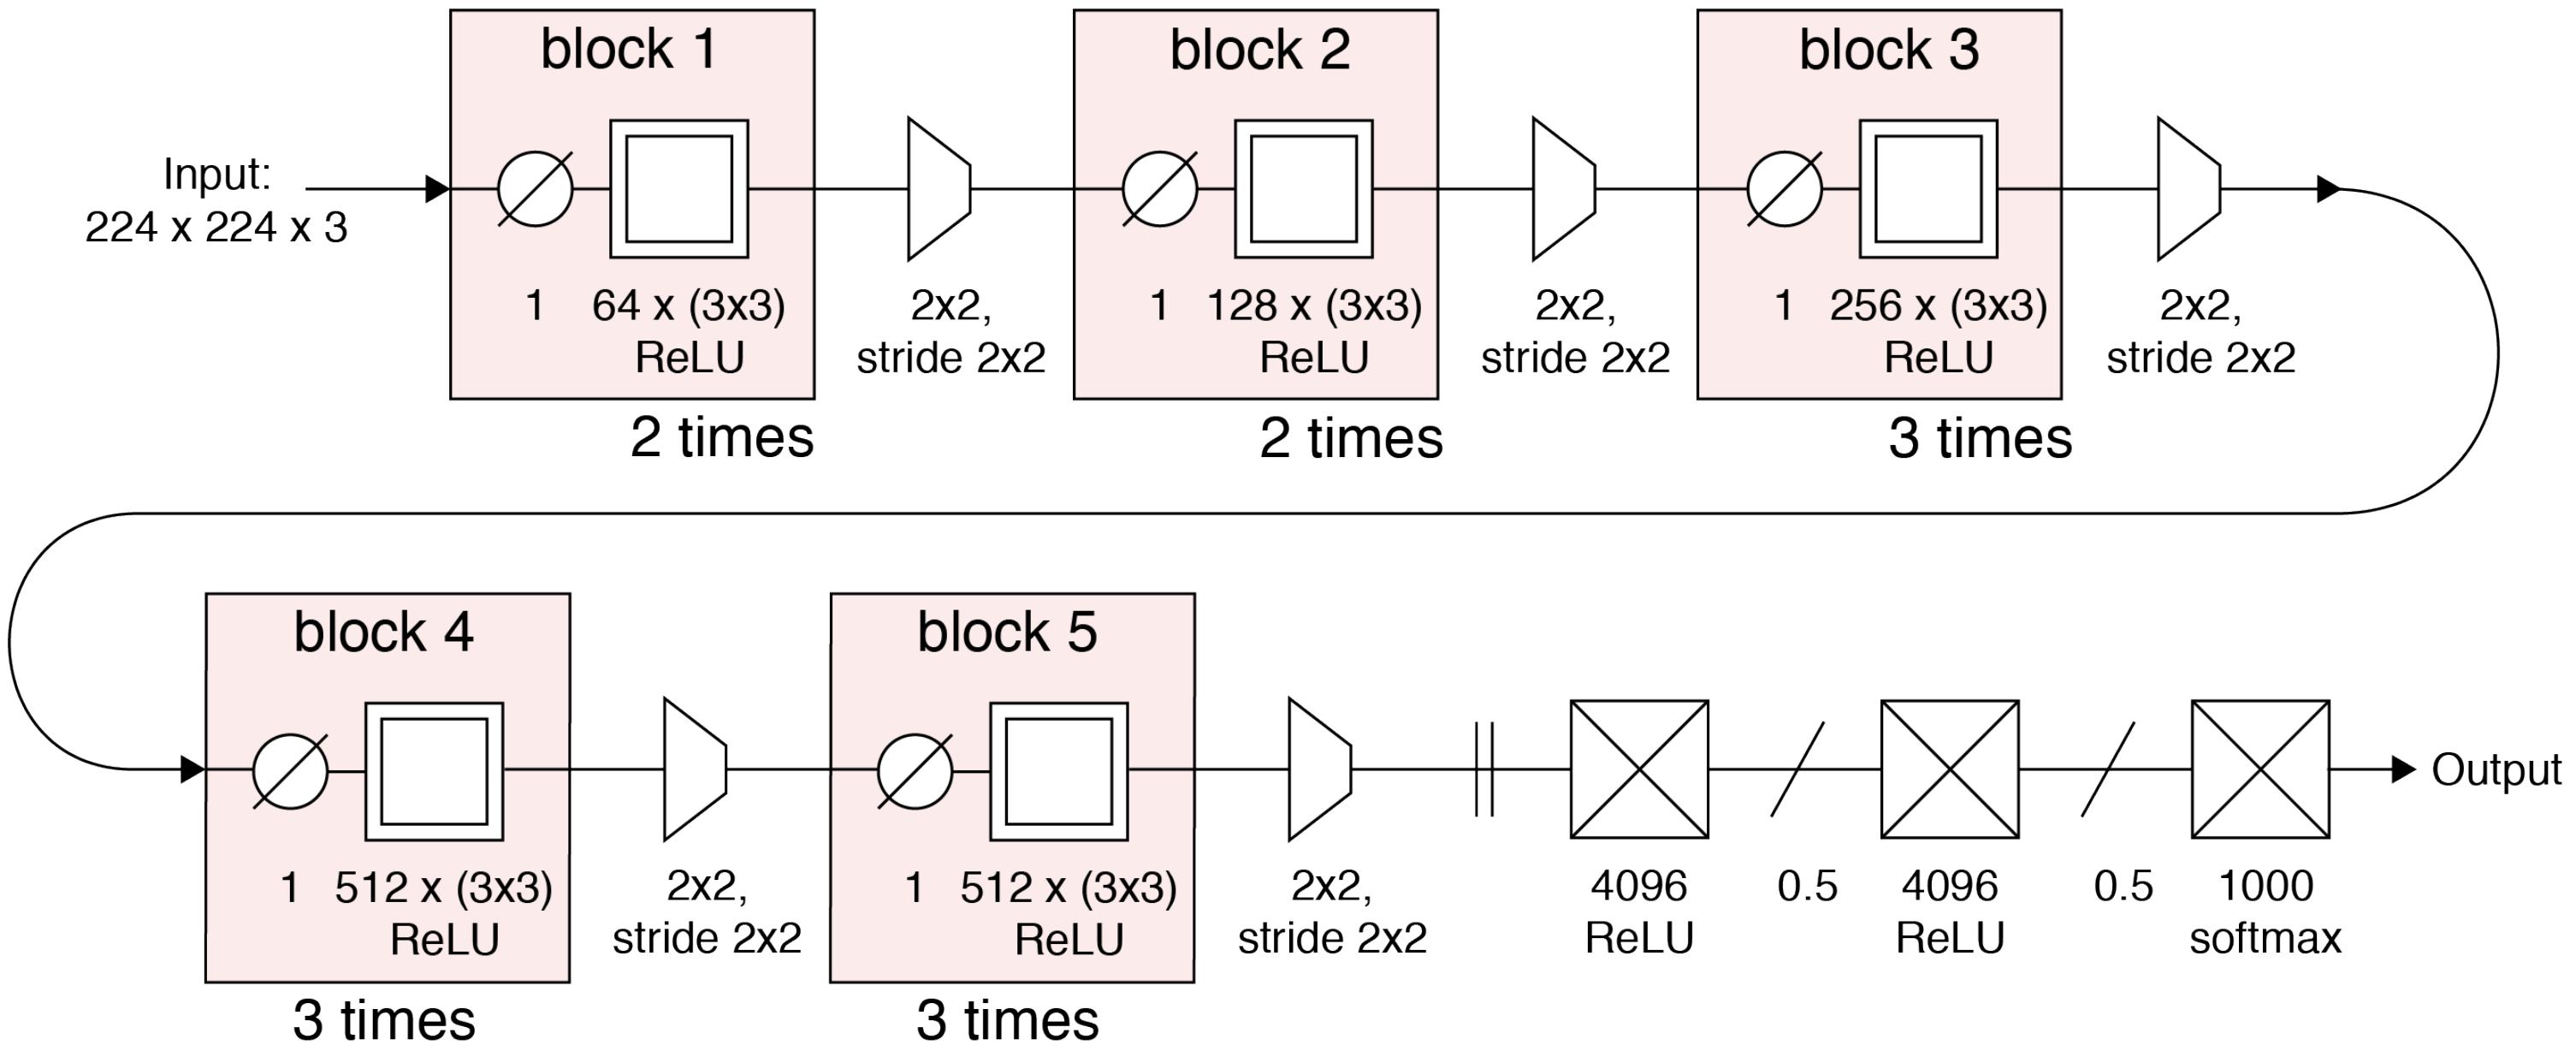
- Image is scaled down at each step,
- Filters work with lower-res versions
- filters can more easily look for larger features



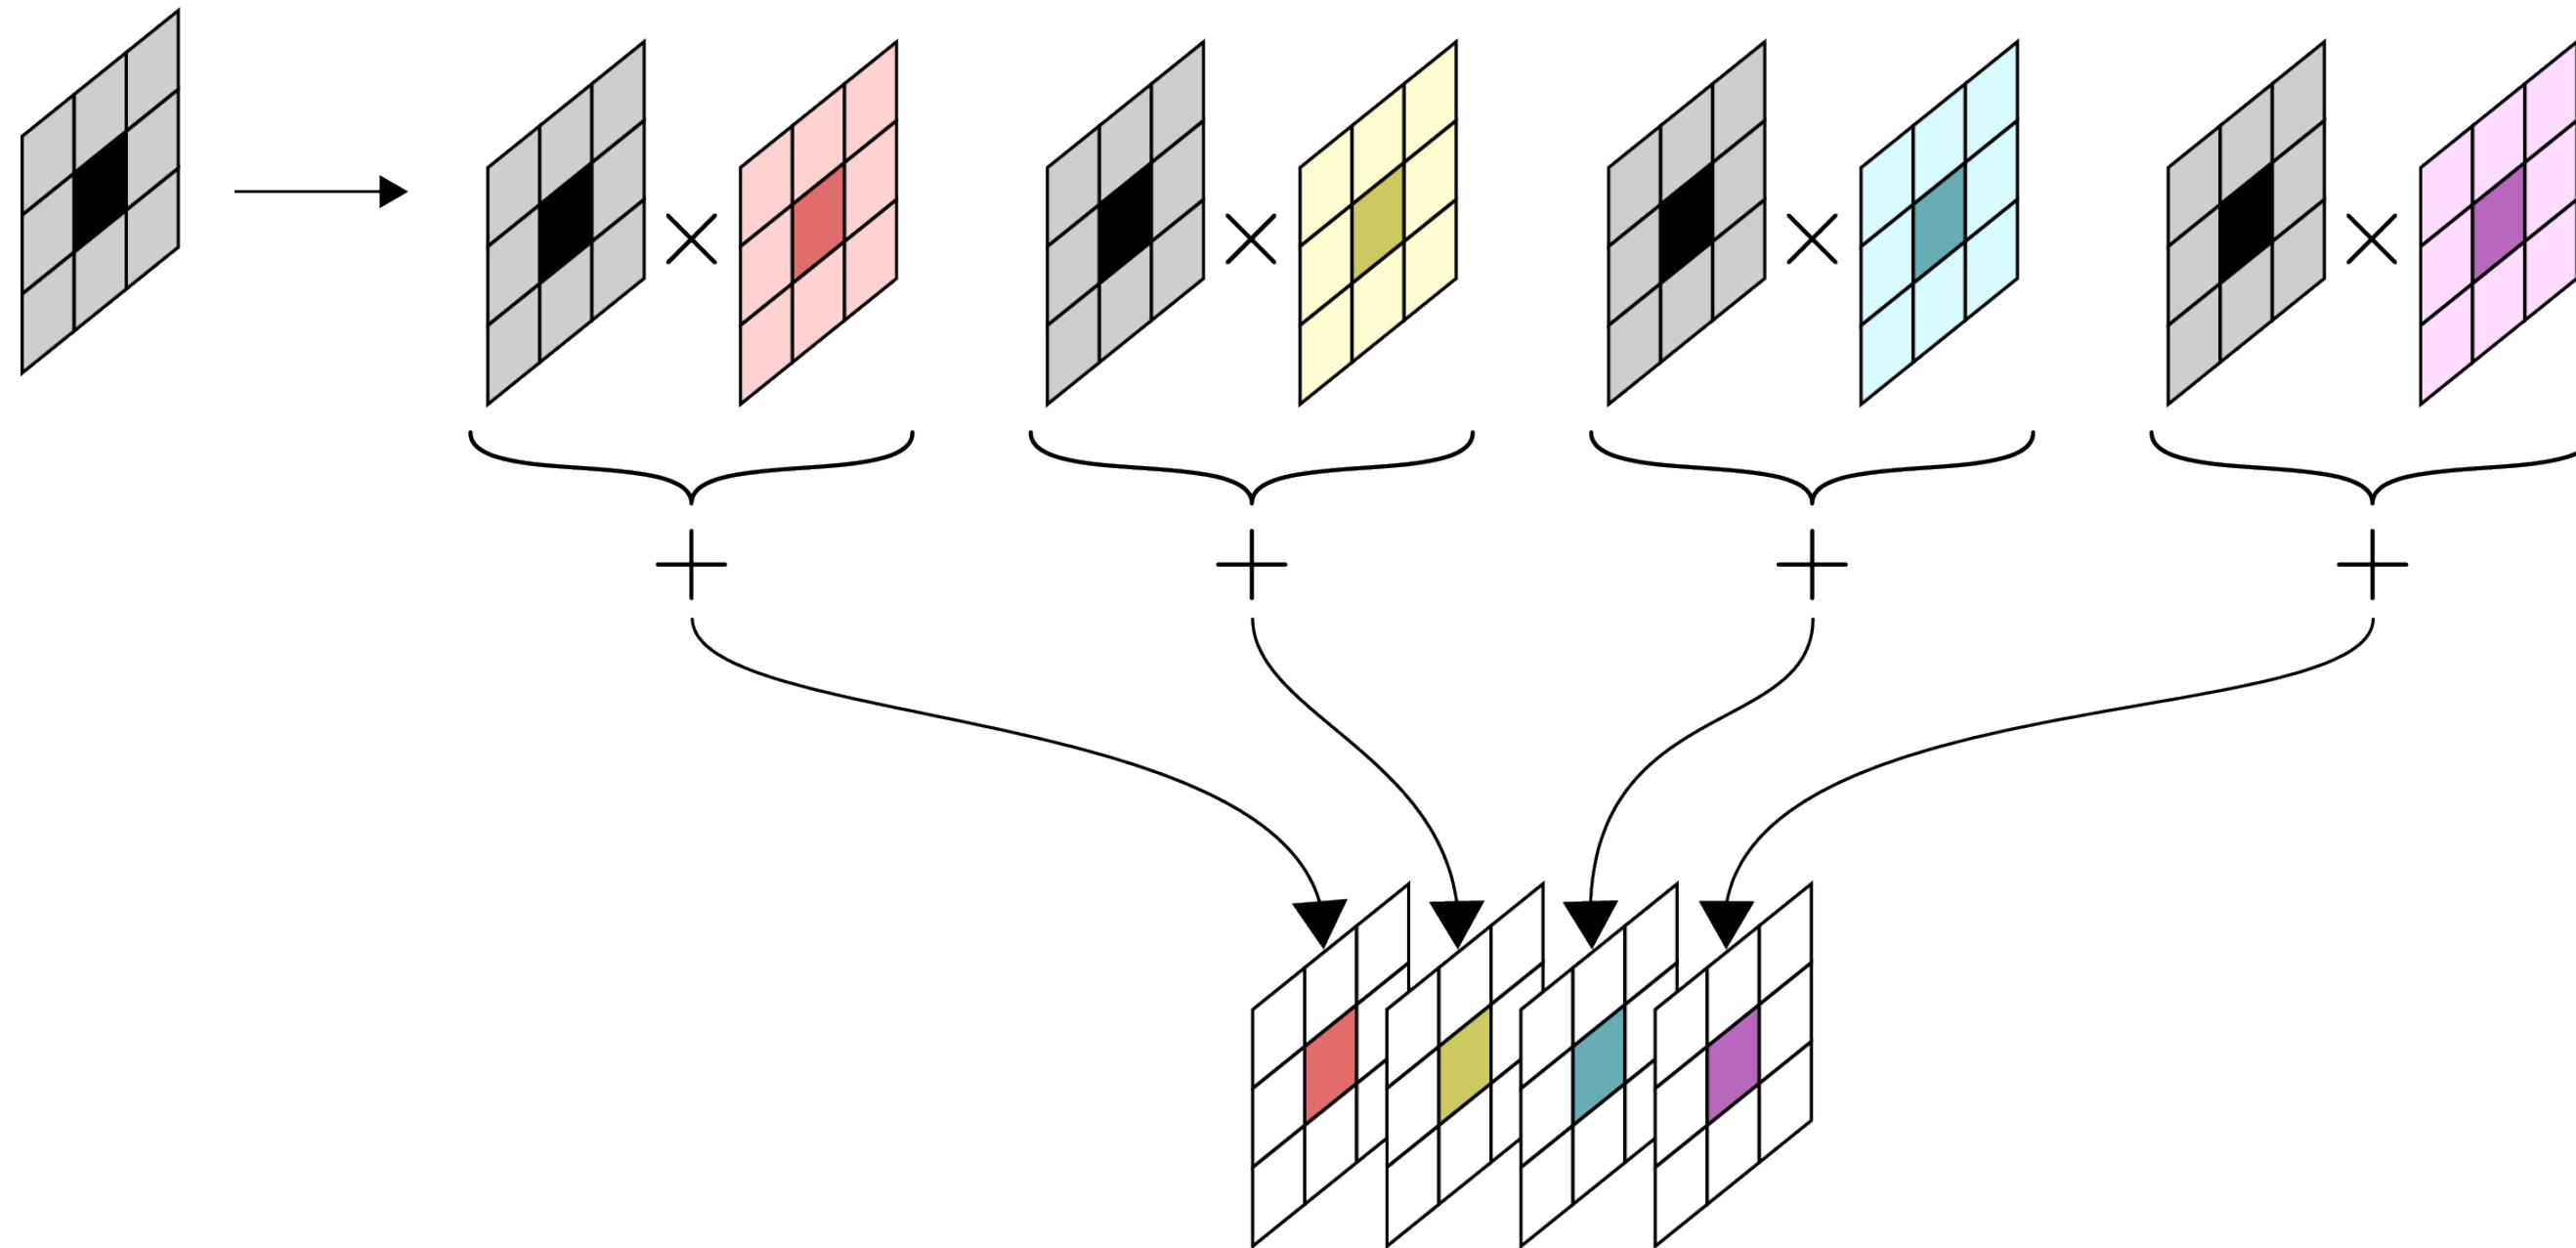
# Layer types schematic



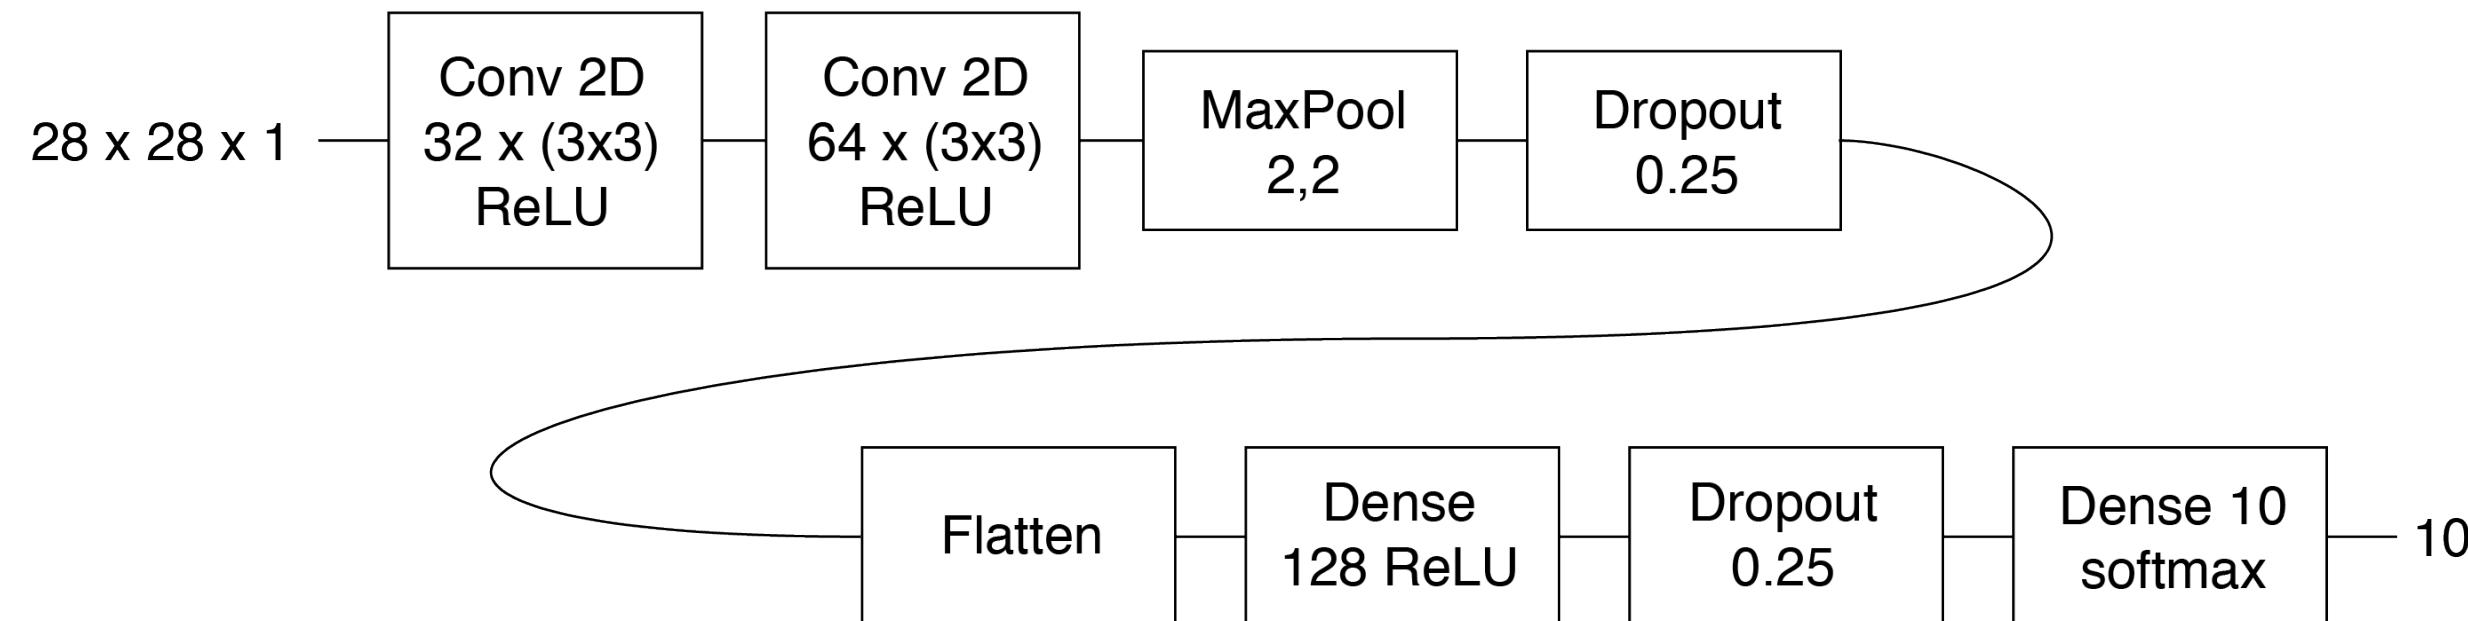
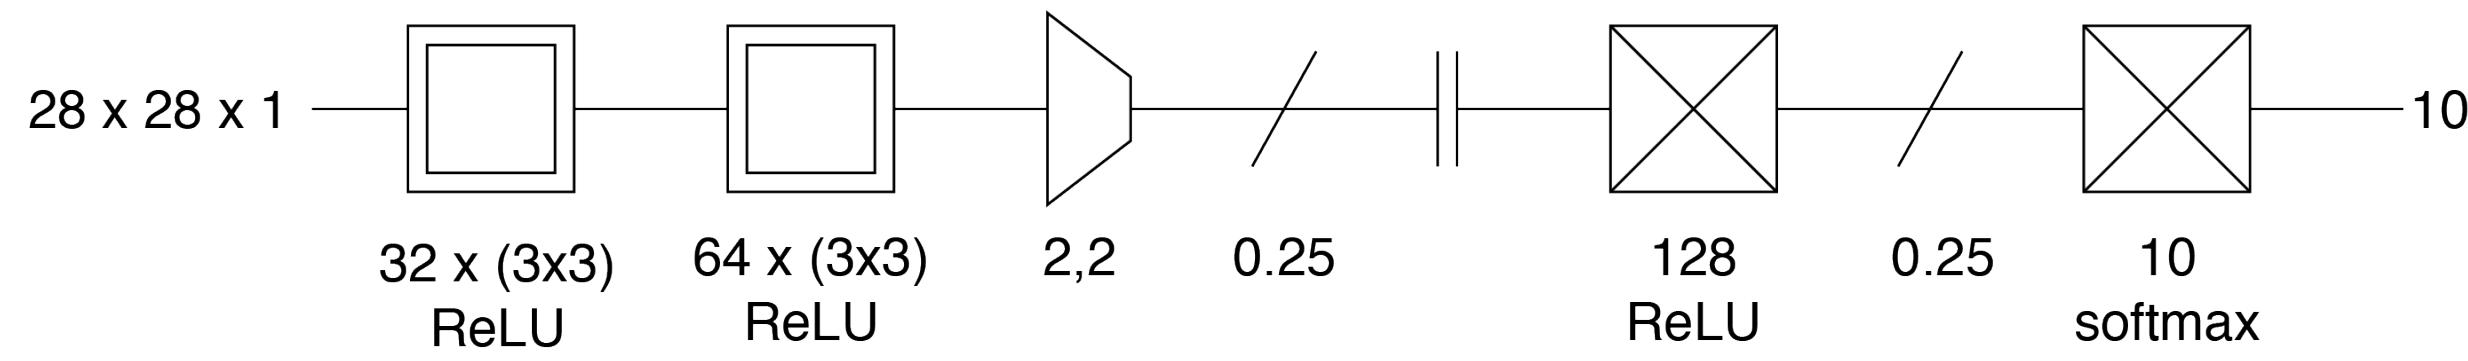
# VGG16 Image Classification Example



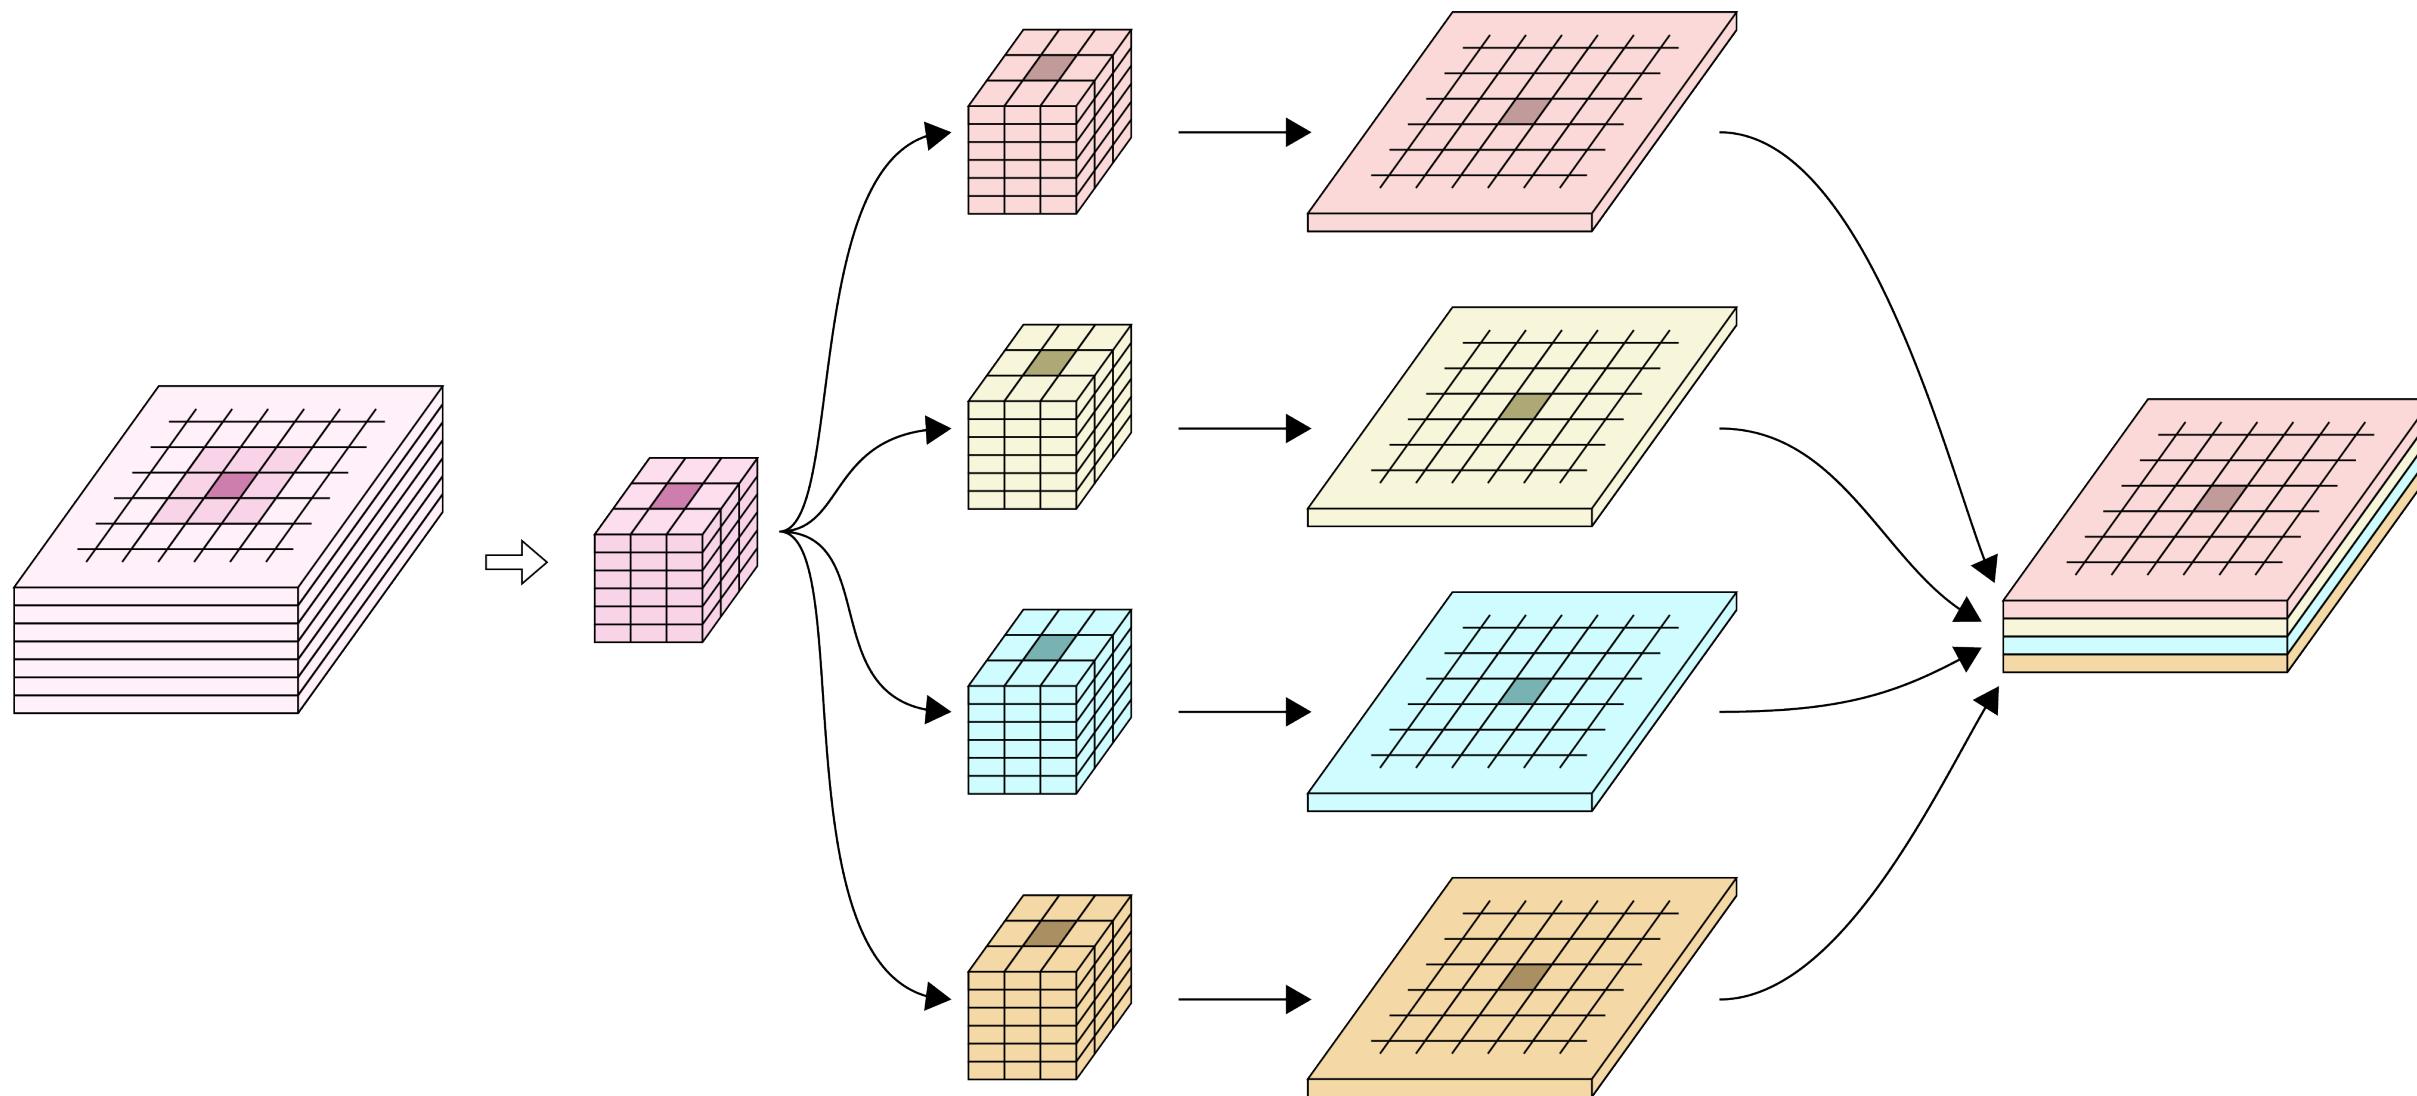
# How Channels work: each color a different feature



# MNIST architecture from Keras Docs



# Channel Arithmetic



input is (say)  $26 \times 26 \times 6$ , so filters MUST have 6 channels and we have 4 new featuremaps:  $\text{Conv2D}(4, (3, 3))$